

APM 598: Homework 1 (2021/02/04)

Ex 1. [Overfitting]

Some data $\{x_i, y_i\}_{i=1\dots N}$ has been generated from an unknown polynomial $P_*(x)$ (see figure 1 and code page 3 to load the data 'data_HW1_ex1.csv' in Python), i.e.

$$y_i = P_*(x_i) + \varepsilon_i$$

where ε_i is some 'white noise' or more precisely $\varepsilon \sim \mathcal{N}(0, \sigma^2)$ with σ^2 is the intensity of the noise (also unknown). The goal is to estimate the polynomial P_* and in particular its degree denoted k_* (i.e. $k_* = \deg(P_*)$).




Figure 1: Data points $\{x_i, y_i\}_{i=1\dots N}$ generated from a polynomial P_* .

- a) For each $k = 0 \dots 12$, estimate the polynomial \hat{P}_k of order k that minimizes the mean-square-error (use *polyfit*):

$$\ell(P) = \frac{1}{N} \sum_{i=1}^N |y_i - P(x_i)|^2,$$

in other words find the polynomial \hat{P}_k satisfying:

$$\hat{P}_k = \underset{P, \deg(P) \leq k}{\operatorname{argmin}} \ell(P).$$

Plot the loss $\ell(\hat{P}_k)$ as a function of k (i.e. plot the points $\{k, \ell(\hat{P}_k)\}_{k=0\dots 12}$).

- b) Split the data-set $\{x_i, y_i\}_{i=1\dots N}$ into training and test sets using (respectively) 80% and 20% of the data. We define two loss functions:

$$\ell_{\text{train}}(P) = \frac{1}{N_{\text{train}}} \sum_{i \text{ in train set}} |y_i - P(x_i)|^2 \quad (1)$$

$$\ell_{\text{test}}(P) = \frac{1}{N_{\text{test}}} \sum_{i \text{ in test set}} |y_i - P(x_i)|^2 \quad (2)$$

where N_{train} and N_{test} denote the number of elements in (respectively) the train and test sets ($N_{\text{train}} + N_{\text{test}} = N$).

For each $k = 0 \dots 12$, estimate the polynomial \tilde{P}_k of order k that minimizes the mean-square-error over the training set:

$$\tilde{P}_k = \underset{P, \deg(P) \leq k}{\operatorname{argmin}} \ell_{\text{train}}(P).$$

and plot the values of the loss functions ℓ_{train} and ℓ_{test} at \tilde{P}_k for all k .

- c) Guess what is the order k_* of the polynomial P_* and give your estimate of its coefficients.

Ex 2. [Gradient descent]

The goal is to implement and test gradient descent methods. We use linear regression as a toy problem using the data set $\{x_i, y_i\}_{i=1\dots N}$ from **Ex 1** and we would like to minimize the following loss function:

$$\ell(a, b) = \frac{1}{N} \sum_{i=1}^N |y_i - (a + bx_i)|^2.$$

- a) Compute the gradient of the loss function $\nabla \ell = (\partial_a \ell, \partial_b \ell)$.
Deduce (numerically) the minimum (a_*, b_*) of ℓ .
- b) Starting from $(a_0, b_0) = (1.8, 1.4)$ and using the constant learning rate $\eta = .05$, implement the gradient descent algorithm:

$$\begin{pmatrix} a_{n+1} \\ b_{n+1} \end{pmatrix} = \begin{pmatrix} a_n \\ b_n \end{pmatrix} - \eta \nabla \ell(a_n, b_n).$$

Estimate the rate of convergence (i.e. how fast does $\|(a_n, b_n) - (a_*, b_*)\|$ converge to zero?).

- c) Implement the momentum and Nesterov algorithm, denoting $\theta_n = (a_n, b_n)$,

$$\begin{array}{ll} \textbf{momentum} & \begin{cases} v_{n+1} = \gamma v_n + \eta \nabla \ell(\theta_n) \\ \theta_{n+1} = \theta_n - v_{n+1} \end{cases} \\ \textbf{Nesterov} & \begin{cases} v_{n+1} = \gamma v_n + \eta \nabla \ell(\theta_n - \gamma v_n) \\ \theta_{n+1} = \theta_n - v_{n+1} \end{cases} \end{array}$$

Estimate the convergence rate for $\gamma = .9$ and $\eta = .05$.

- extra) Test other gradient descent methods (e.g. Adam, AdaGrad, RMSProp) using class defined in either Pytorch or TensorFlow.

Ex 3. [Classification]

The goal is to implement a linear classifier for the data-set MNIST-fashion (see also code page 3).

- Adapt the linear classifier for the MNIST data-set to the new data-set.
- Train the model using 40 epochs with learning rate $\eta = 10^{-4}$ and the gradient descent method of your choices.
Plot the evolution of the loss at each epoch.
- Draw the 'template' of each class.

```
#####
##      Exercise 1      (Python)      ##
#####

import matplotlib.pyplot as plt
import numpy as np
# get the data
data = np.loadtxt('data_HW1_ex1.csv',delimiter=',')
x,y = data[:,0],data[:,1]
# plot
plt.figure(1)
plt.plot(x,y,'o')
plt.xlabel(r'$x$')
plt.ylabel(r'$y$')
plt.show()

#####
##      Exercise 3      (Python)      ##
#####

from torchvision import datasets
import matplotlib.pyplot as plt
# Download and load
data_collection = datasets.FashionMNIST('data_fashionMNIST', train=True, download=True)
# Illustration
label_fashion = dict([(0, 'T-shirt'), (1, 'trouser'), (2, 'pullover'), (3, 'dress'), (4, 'coat'),
                      (5, 'sandal'), (6, 'shirt'), (7, 'sneaker'), (8, 'bag'), (9, 'boot')])
X,y = data_collection.__getitem__(123)
plt.figure(1);plt.clf()
plt.imshow(X)
plt.title("Example of image with label "+label_fashion[y.item()])
plt.show()
```